

CoStor, a peer-to-peer distributed backup solution

Robert Phipps

4th Year Project Report
Computer Science
School of Informatics
University of Edinburgh

2020

Abstract

This is an example of `infthesis` style. The file `skeleton.tex` generates this document and can be used to get a “skeleton” for your thesis. The abstract should summarise your report and fit in the space on the first page. You may, of course, use any other software to write your report, as long as you follow the same style. That means: producing a title page as given here, and including a table of contents and bibliography.

Acknowledgements

This project makes extensive use and builds on top of concepts explored in this paper from Paul Anderson and Le Zhang: Fast and secure laptop backups with encrypted de-duplication [1]

Table of Contents

1	Solution overview	7
1.1	Goals of CoStor	8
1.2	Limitation of scope for the purposes of this project	9
1.3	Exploration of existing solutions	10
1.3.1	Syncthing	10
1.3.2	UrBackup	10
1.3.3	Hermes	10
1.3.4	Bacula	11
1.3.5	Amanda	11
1.3.6	GlusterFS	11
1.4	Deployment topology	12
1.5	Specimen use case	12
2	Client implementation	15
2.1	Summary	15
2.2	The backup process	15
2.3	Large file upload API	16
2.4	Management and monitoring	17
2.5	SQL database and other backup plugins	17
2.6	Filesystem snapshots and integrity protection	17
2.7	The HashTreeMaker	18
3	Server and datastore implementation	19
3.1	Current development status	19
3.2	Overview	19
3.3	High-level data architecture	20
3.4	Server metadata database and filestore	20
3.5	Management UI	21
3.5.1	Monitoring	21
3.5.2	Data browser	21
3.6	Restoring backups	21
3.6.1	Direct (local) restore	21
3.6.2	Remote (disaster recovery) restore	22
3.6.3	Local (disaster recovery with new server) restore	22
3.7	Automated maintenance tasks	22

4	Multi-site replication	25
4.1	Data distribution and sharding	25
4.2	Replication management database	25
4.3	Securely synchronising data	26
4.4	Efficiently transferring data	27
5	Testing and evaluation	29
5.1	Being a reliable backup solution	29
5.2	Being easy to use and manage	29
	Bibliography	31

Chapter 1

Solution overview

CoStor is designed as a turnkey solution to the problem of maintaining reliable system backups within an SME with multiple sites, such as a confederation of schools. Instead of using expensive and bandwidth intensive cloud storage services for offsite backup, CoStor is designed to hold a complete local backup on-site within the CoStor server, as well as automatically replicating backup data across a group of federated instances of the server software, ensuring that there is always at least two redundant copies of the backup datastore in two different physical locations.

To simplify networking requirements for deployment, all communication between clients and servers, both locally and between sites, makes use of standard HTTPS requests. This negates the need for complex multi-site VPNs, and simply requires a single TCP port to be forwarded to the server from the internet.

The backup datastore's metadata and directory structures are maintained inside an SQL database and can only be modified over the RESTful HTTP API, reducing attack surface compared to making use of more traditional file transfer methods like FTP, NFS or SSHFS.

All management is completed through a simple web UI, where clients can be configured for one-touch deployment and subsequently monitored, as well as where users can browse the directory trees for each backup "snapshot" in the event that a file needs to be recovered. A full backup restore can be completed by requesting a restoration archive, whereby the server will build a complete archive of a snapshot, pulling data from its local datastore, or from other sites in the event of a remote restore.

This system makes use of Django for server-side components, a Python web framework with fantastic ORM and enforcement of best-practices.

1.1 Goals of CoStor

Given the target organisations for CoStor, there are some specific goals that need to be targeted during development.

- **Reliable backups**

As should be very obvious, being a backup solution, CoStor needs to be able to reliably manage and maintain backups for a network. This includes protections such as an "append-only" API for backup clients, validation of uploaded data, and mitigations against the most common reasons a backup may be called upon such as accidental deletion, user errors, hardware failure and ransomware style attacks.

- **Simple restores**

As this system is targeted at small organisations which may not have their own full-time IT support staff, restoring from backups should be straightforward for an end-user. This is achieved through the use of a self-service web UI.

- **Robust security and audit logs**

Backups almost always contain confidential information, so CoStor needs to be able to manage permissions on a granular user-by-user basis. It also needs to include audit logging for all operations on the system. Any offsite storage and "data in flight" needs to be strongly encrypted to protect confidentiality.

- **Low maintenance**

Systems tend to be forgotten about, and in the case of backups, often you only notice something hasn't been working once you need to restore something¹, so CoStor needs to include built in maintenance and scheduled self-tests to ensure that the system is ready when the user needs it most.

- **Simple deployment**

Again, CoStor needs to be deployable by inexperienced IT support staff without prior knowledge of network filesystems, command line interfaces or web development. This can be achieved by making use of clever packaging and deployment strategies such as Docker for server components and zero-touch installation scripts for backup clients. By making use of standard and well understood protocols such as HTTP(S) for communication between components, compatibility with most network architectures should be maintained, without the requirement for complex network share configurations, multi-site VPNs and authentication systems.

- **Centralised management and monitoring**

¹GitLab found this to their cost in 2017 after discovering their backups hadn't been running for some time: <https://techcrunch.com/2017/02/01/gitlab-suffers-major-backup-failure-after-data-deletion-incident/>

As this is designed to be deployed over a large number of client PCs, ensuring configuration is correct could be challenging. As such, CoStor will include configuration of backup clients from the server management panel, and all clients will pull down configuration automatically from this central repository. Backup logs will also be pushed to the server so that an administrator can monitor their entire estate from a single place.

- **Distributed and fault-tolerant file stores**

Leaving the best to last, CoStor's standout feature will be that backups can be automatically replicated between federated instances of the CoStor server, over a zero-configuration HTTPS link. The system should be able to recover from the loss of a server without any data loss, and allow restoration of data originating from any site from any of the remaining instances of CoStor within the network.

1.2 Limitation of scope for the purposes of this project

As this project is to be completed by an individual over the course of a single academic year, the scope does unfortunately have to be limited somewhat. The primary goal is to complete a "Beta" release of the server application, with cross-site replication, web UI and ability to restore data, along with a barebones backup client to allow the system to be demonstrated.

Managing filesystem metadata and ensuring consistent snapshots is a considerable problem in its own right.

1.3 Exploration of existing solutions

There are many packages in existence which incorporate a subset of the features targeted by this project, however most either focus on the synchronisation features with some limited support for file history, and no robust backup capability, whereas others rely on cloud storage infrastructure as the backend, distributed datastore which either necessitates the use of a commercial provider such as AWS S3² or BackBlaze B2. Both of these greatly increase cost and introduce a reliance on a third party.

A number of existing products have been selected as they offer the closest functionality to that targeted by CoStor, and are explored here:

1.3.1 Syncthing

Syncthing is a service targeted at consumers who want a self-hosted alternative to commercial cloud storage and synchronisation services such as Dropbox, Google Drive and OneDrive. The agent software can be configured to sync any file changes between two or more devices, allowing for access anywhere, with a form of georeplication. It also requires fairly minimal network configuration, just needing a pair of ports to be opened on at least one of the nodes to allow discovery of other agents.

A very large distinction has to be made in the fact that "Syncthing is a continuous file synchronization program"[2], in that it isn't designed to create restorable snapshots of your data, and therefore is completely unsuitable for robust *backup* of important information.

More information is available at <https://syncthing.net> [2]

1.3.2 UrBackup

UrBackup is closer to CoStor in its goals, as is specifically built to be used as a backup system. It can auto-discover agents on the network, and begin incremental backups to its server software. It also supports full image backups of NTFS formatted drives with bare metal restore. UrBackup has many of the features targeted by CoStor, including a simple web interface for management, however it does not include any built-in support for georeplication.

More information is available at <https://www.urbackup.org/index.html> [3]

1.3.3 Hermes

Hermes is an "open-source redundant distributed storage network"[4]. Although it doesn't come pre-packaged with components allowing it to be used as a turnkey backup

²Amazon Web Services' bucket storage solution

system, it is worth exploring as it does specifically target the geo-replication features for the purposes of backup that CoStor is looking to integrate. It promises fast, encrypted and seamless replication and sharding of data across nodes in the network, making use of LZMA compression to increase performance.

This would appear to be a very promising option to integrate as the backend storage for CoStor, however it looks like the project is very much stale, with the last commits being made in late 2014. As such, its codebase is somewhat limited in utility, given its lack of ongoing maintenance. It is also written in Go, with limited documentation, which is not a language that I am familiar enough with to begin work on reviving.

More information is available at <https://github.com/Hermes/hermes> [4]

1.3.4 Bacula

Bacula is an open-source and very mature backup framework, with tools to allow a multitude of network configurations. Unfortunately its flexibility does result in the software being complex to configure. CoStor is targeting small organisations with limited in-house IT support capacity, so this would likely be too complex to deploy without the assistance of external contractors. Bacula is also available in an "enterprise" edition[?], which includes support as part of the subscription cost, however this version is both closed-source and not inconsiderably expensive.

More information is available at <https://www.bacula.org> [?]

1.3.5 Amanda

Amanda is another backup-specific solution, again with a "community edition" being accompanied by a commercially supported "enterprise" version of the software. Like the other systems investigated, it targets backup to local storage, NAS drives and traditional tape backup libraries, so would not be suitable for the use case targeted by CoStor.

More information is available at <https://www.zmanda.com/amanda-community-edition.html> [?]

1.3.6 GlusterFS

The Gluster filesystem is a "software scalable network filesystem"[?], which allows systems administrators to build complex filesystems across multiple physical machines using commodity hardware. Gluster is targeted at engineers designing private cloud environments which need large volumes of storage accessible from many devices, and has a huge number of configuration options to allow, among other things, redundant replicated storage across machines. It also has support for filesystem filters that can offer encryption at a filesystem level before data is transmitted between nodes.

Unfortunately for CoStor's usecase, this system is built on the assumption that the Gluster volumes are configured at the point of deployment, by the same user on all sites. It also is primarily managed through a command line interface, with limited Python bindings, and is targeted at an architecture where all nodes have rights to see data within the volume. Finally, it is based around a master-slave architecture, which is unsuitable if CoStor were to be using a shared volume for all nodes. It would be possible to wrangle a Gluster configuration to match what is required by CoStor, where each site has its own volume, with its own georeplication session, however this would be very difficult to automatically provision from the CoStor web interface, and should be mostly unnecessary given the data structure used for backups.

More information is available at <https://www.gluster.org> [?]

1.4 Deployment topology

The basic topology of a CoStor network would be as follows:

- **Clients** (many):

The CoStor client software is installed on any systems which are to be backed up. It communicates over the local network to the site's local instance of the CoStor server.

- **Servers** (one per site/network):

There should be one instance of CoStor server on the internal private network of any site that has clients to be backed up. There could be multiple instances running in one local network space, but they would operate as separate "sites" within the software. The servers are the primary file store for a site, and manage replication between federated instances on other sites.

- **Federated servers** (> 2 across multiple locations):

These are additional instances of the server application, and can be on separate networks and in different physical locations. These communicate over the internet through a replication API to keep a redundant copy of all data from all servers. They can restore data from any site within the group in the event that a site's server fails, assuming the target site's encryption key is provided.

TODO: Insert diagram of topology, and describe how data is distributed across the members of the CoStor network.

1.5 Specimen use case

The initial inspiration for this project was a use case within a confederation of small, state controlled schools. Each school currently spends a large amount of money on a

third-party managed cloud backup solution - in the region of £500 per year for only 20GB of backed up data. This meant that they were unable to backup considerable portions of their data, leaving some of their media based teaching resources vulnerable.

The purpose of the confederation is to allow the members to share resources and purchasing power for resources that can be beneficial to all of them.

CoStor could be utilised by these schools to allow them to geo-replicate their backups among themselves, on commodity server hardware, allowing them to no longer be reliant on expensive cloud backup solutions. Each site would have a single CoStor server appliance with their file servers, databases, Microsoft Active Directory databases and any other critical data being backed up by instances of CoStor client, allowing backups to complete rapidly across the high speed LAN, with replication to other schools' CoStor servers being carried out overnight when the internet connection is quiet.

The system should be as close to turnkey and zero-maintenance as possible, allowing it to be set up by their IT support contractors, and left to work for years at a time without much interference.

Chapter 2

Client implementation

2.1 Summary

CoStor clients are designed to be managed centrally through the CoStor server web interface, therefore they require no GUI of their own, just a way to set the target server instance, required directory tree to backup (the backup root) and their authentication token. These are set using a `.YAML` file, which is easy to push to clients with standard network management tools such as Microsoft Group Policy and Logon Scripts.

The client is built around Python 3, and PonyORM working in parallel with an SQLite database. This means there are very few dependencies, once the Python project is bundled, and allows cross-platform development with a shared codebase.

2.2 The backup process

The backup process works as follows:

1. **Build directory structure hashes:** The client uses the `HashTreeMaker` class to build a simplified version of the directory tree being backed up, traversing through the tree bottom-up, and building objects for every item it encounters, each item identified by a hash based off of its attributes and the hashes of all its children. File data is also hashed and stored with a path to the file, to be used for backup later.
2. **Local metadata database update:** The client keeps a local database of all backup "snapshots" as trees of objects, created by the hasher. It checks for an existing snapshot, and uses this to identify any subtrees within the backup root which have changed since the last backup run. It then writes any changes to the database as part of a new snapshot.
3. **Authenticate with server:** A quick test is made to ensure that the client's credentials are valid and that the server is accessible.

4. **Push changed metadata objects:** The server is queried to see if it has seen any objects with matching IDs (hashes), and any objects that are missing are pushed as a single JSON object.
5. **File "primes" are pushed to the server:** Again, the server is queried to check for existing copies of any of the files that are included in this backup snapshot, by hash. For any files that are missing, the client uses the path included with the object to find a copy of that file, and pushes it to the server, in multiple chunks.

2.3 Large file upload API

To allow large files to be uploaded over standard HTTPS connections, with the ability to recover from connection issues, files are split into (currently 100MB) chunks, which are individually hashed, and then pushed in sequence to the server.

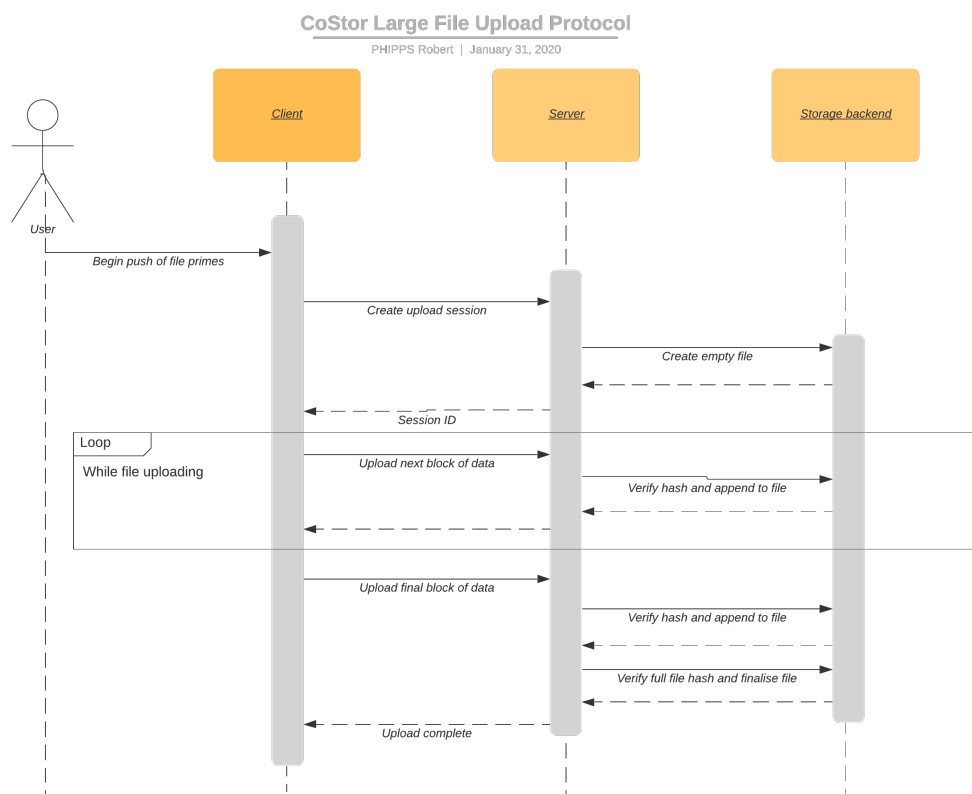


Figure 2.1: UML sequence diagram for large file upload

Integrity is ensured by having the client create an "upload session" before pushing chunks of files. This session includes the file hash and the number of expected parts. Each part is also uploaded along with its hash and sequence number, which is verified before the data is appended to the file in the server's filesystem. Once the final part has been received, the server checks the completed file against the hash given in the upload session before marking the file as complete and closing the session.

Chunks can be uploaded over any period of time, and chunks for one file/session can be uploaded interspersed with chunks from other files, the only constraint is that chunks for a session are received in correct sequence order.

Should a chunk fail its upload, due to a verification error or network instability, the client is able to retry this chunk once network has been restored, allowing uploads to be "resumed", at least to the resolution of the chunk size.

2.4 Management and monitoring

Telemetry and configuration synchronisation is currently not implemented

On each run, the client checks in with the server to pull down a configuration update, which includes chunk size, backup schedules, backup root, server version, and expiry time for local metadata information. This allows it to ensure it is configured correctly, without needing local intervention from the administrator.

It also pushes up telemetry from backups including backup results and network speed statistics, as well as free disk space and local metadata database size.

2.5 SQL database and other backup plugins

Plugins are currently not implemented

Included in the client are plugins to allow the client to automatically dump and backup common applications such as Microsoft SQL Server. These plugins are simple Python modules installed in the CoStor client install directory. Simple command line scripts can also be defined from the CoStor server interface and pulled down with the client config package.

Current targeted software includes:

- Microsoft SQL Server
- MySQL server (Windows/Linux)
- Microsoft Active Directory Server¹

2.6 Filesystem snapshots and integrity protection

The below has not yet been implemented

¹<https://docs.microsoft.com/en-us/windows/win32/ad/backing-up-an-active-directory-server>

To ensure that files cannot change during the backup process, it is necessary to create some form of snapshot of the filesystem while the backup is running. This is currently being investigated and will likely make use of LVM, Microsoft Shadow Copies and Time Machine on Linux, Windows and MacOS respectively.

2.7 The HashTreeMaker

This class is the engine behind the creation of the filesystem metadata snapshot, and builds simple objects that can be stored in the relational databases used by CoStor.

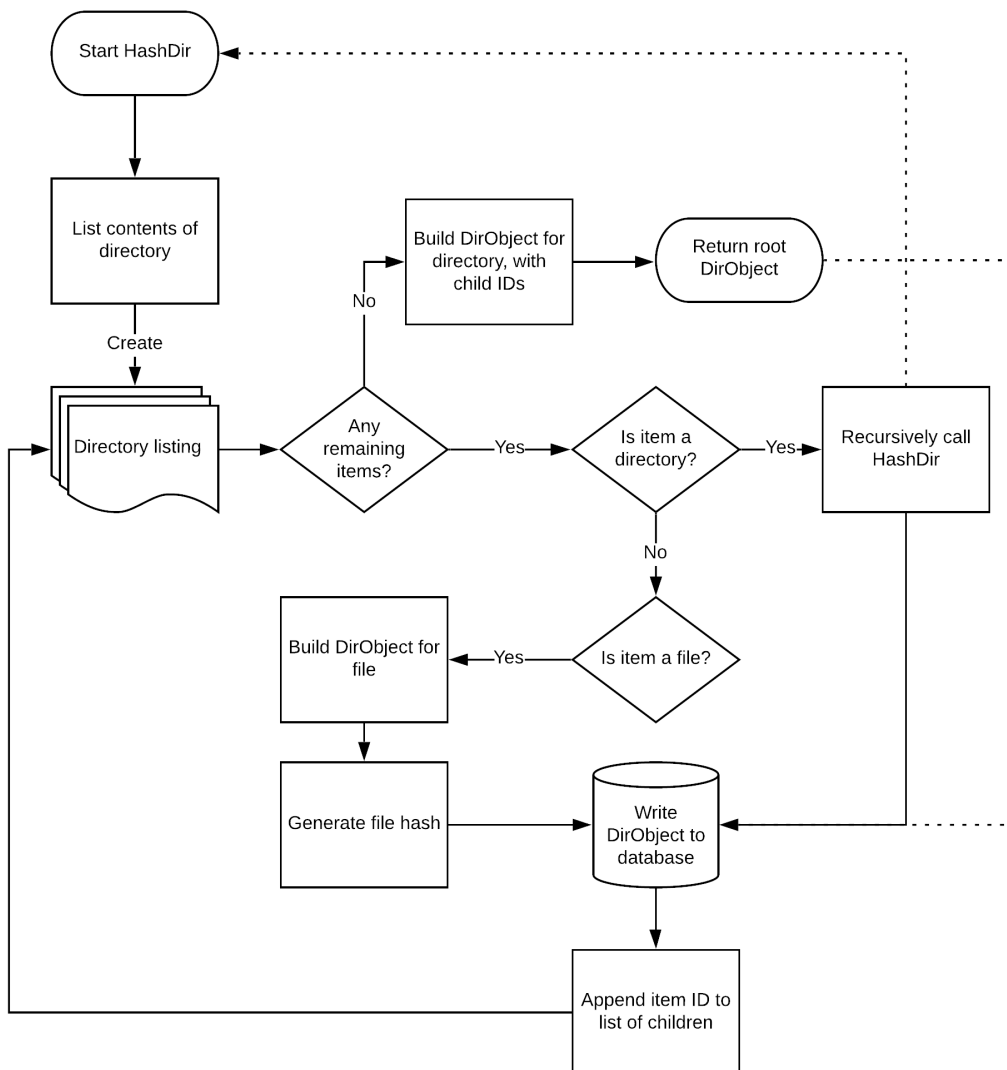


Figure 2.2: Process for creating DB representation of directory tree

This currently uses SHA1, which is not secure. TODO: move to BLAKE2

Chapter 3

Server and datastore implementation

3.1 Current development status

As of late January 2020, the CoStor server application is very minimal, simply accepting backups into its datastore, and all user interface is through the built in Django Admin pages API. As standard model definitions are used, along with the Django REST Framework¹, building API endpoints to expose this data to a custom front-end application should be incredibly straightforward. Work to enable long-running and complex tasks is ongoing.

3.2 Overview

CoStor server is a Django based web application, that can be deployed within Docker for easy installation and upgrades. It not only provides the backup API endpoints used by the Client software to actually push snapshot data to the server's storage, but also a web UI for management and monitoring, as well as frameworks to allow replication of backup data between instances of the server over a standard HTTPS connection.

To allow larger tasks (such as backup archive generation) to run asynchronously of the HTTP requests to the web server, a task broker such as Celery² is used.

¹<https://www.django-rest-framework.org>

²<http://www.celeryproject.org>

3.3 High-level data architecture

To allow all machinery to be as generalised as possible, all possible file system objects are stored as generic `Object` objects. These can represent files, directories or symlinks. Each object holds a globally unique ID, path information, the hash of the object's contents, child and parent relations, the path and link to a `DbFile` object or "prime" in the case that the object is a file, with the prime representing the binary data of that file on the CoStor server's filesystem.

These objects are then related to their associated "snapshot" or backup session, backup root directory and the agent (or CoStor client identity) that the backup originated from.

3.4 Server metadata database and filestore

All data on the CoStor server is stored within the context of the Django ORM, to allow easy access to all objects using the powerful Object APIs provided by the ORM. The backup data is stored in the following database models:

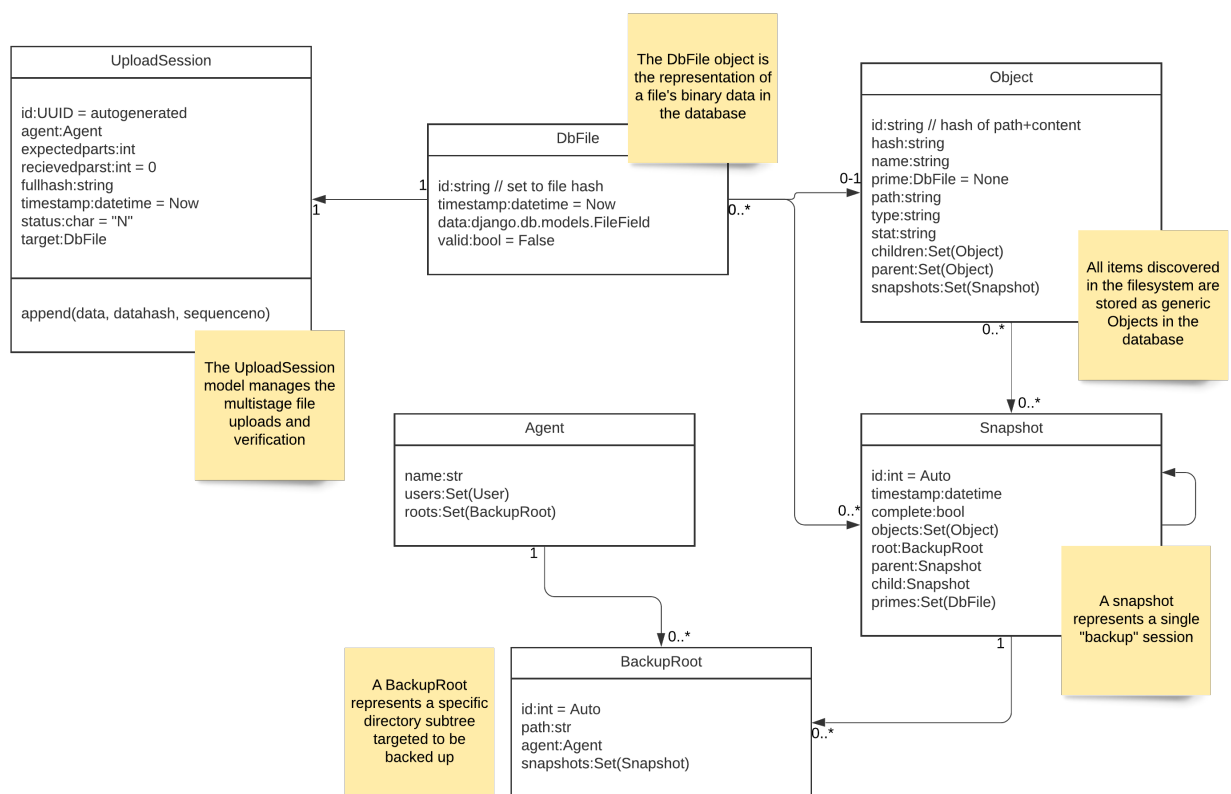


Figure 3.1: UML of Django ORM models used to manage backup data

File data is stored wrapped within a `djando.orm.models.FileField` attached to the `DbFile` object and stored on the filesystem of the CoStor server at a location defined

by the config file, with the filename being the object's ID.

3.5 Management UI

This user interface provides access to all day to day configuration tasks, shows telemetry for all clients and replication status and allows an administrator to browse the directory tree of backup snapshots for any agent enrolled against the server.

3.5.1 Monitoring

Below is a wireframe of the monitoring view for CoStor clients enrolled against the server. This view allows an administrator to check at a glance whether agents have checked in recently, total and individual disk usage and any recent errors. It also shows the result of recent automatic maintenance tasks.

3.5.2 Data browser

Mockup of data browser UI

3.6 Restoring backups

Restoration has not yet been fully implemented

There are two main procedures for restoring backups from CoStor, depending on whether the data is stored on the local server (a direct restore) or whether data is being restored from geo-replicated copies of the database, in the event the site's own instance has failed.

3.6.1 Direct (local) restore

These are the most straightforward of restores, the administrator simply has to find a directory or object within the data browser interface for an agent and snapshot within the web UI for their site's local instance of CoStor server, before requesting a restore package be generated. CoStor then re-builds the original directory tree selected from database objects and creates an archive of the backup which can be downloaded and restored manually.

3.6.2 Remote (disaster recovery) restore

In the event that the instance of CoStor has failed for the site that is being backed up, there are additional steps.

1. As all metadata for replicated sites is encrypted, a valid master decryption key must be provided to the server being used to restore a backup. This is checked against the master federation database, which is replicated to all servers within the group.
2. The server then identifies which CoStor servers contain copies of the metadata databases for the requested site (using the global master federation database), and requests that metadata be replicated to itself if the data is not already available locally.
3. Finally, the server decrypts the replicated metadata for the requested agent and snapshot, allowing the administrator to generate a restoration archive using the same method for a local restore: browsing the directory tree and selecting required sub-trees to be restored. File "primes" or binary data blocks are decrypted on the fly before being written to the restoration archive.

3.6.3 Local (disaster recovery with new server) restore

In the event that the hardware running CoStor has to be replaced, it is possible to replicate all data for that site back to the new server from the replicated copies across other sites.

1. The new instance should be configured with the same ID, encryption key and access token as the original server, to re-enroll it into the network. The master federation database will then be replicated to this new server.
2. The new (replacement) server can then request that all replicated data from its site is returned to its local storage, and will decrypt each package as it arrives to rebuild the server's original state.
3. Once all data has been restored to the server, the restoration continue as outlined in the direct (local) restore procedure.

3.7 Automated maintenance tasks

To ensure that the CoStor server can run for extended periods of time without requiring regular maintenance, a series of day-to-day tasks are automated. These include:

- **Purging of old data:** CoStor will automatically purge old snapshots and orphaned metadata entries once the retention window set by the administrator has expired.

- **Disk space monitoring:** Should the disk space become limited on the host, an email will be sent to the administrator to notify them that this is the case before backups fail.
- **Hard drive SMART testing:** CoStor can be configured to automatically run periodic SMART tests on hard drives within the server, and notify administrators in the event of an anticipated failure. This very simply runs the `smartctl`³ command and parses the command line response.
- **Replication configuration updates:** CoStor is clever enough to automatically change replication targets, should any of its currently selected targets become unavailable for whatever reason (such as lack of disk space or hardware failure).
- **Offline client notifications:** If an instance of the client hasn't completed a backup in a defined period of time, the administrator will be notified to allow them to investigate.

These tasks are run on a fixed schedule, prior to replication tasks running, in a Celery task runner.

³<https://linux.die.net/man/8/smartctl>

Chapter 4

Multi-site replication

The mechanics for this are currently still being prototyped.

The data to manage multi-site restores is stored in a globally replicated database, which contains information such as the replication status of all snapshot directory trees and file "prime" objects, anonymised agent ID to site mappings, and synchronisation job objects to track when and to where objects are replicated.

Each site manages its own replications, and responds to replication requests from other servers within the network. All objects are identified globally by the hash of their contents, and database objects are replicated as encrypted JSON documents. File "primes" are transferred using the large file upload protocol used by the clients (outlined in Figure 2.1)

The overall goal of the replication system is to ensure there are at least two redundant, offsite copies of all data. This is performed by selecting replication targets randomly, weighted by free disk space, and by incrementally scheduling the replication tasks during a configured maintenance window (normally overnight when the network is quiet).

4.1 Data distribution and sharding

Each server will choose two sites to distribute its data across to. This means that each instance of CoStor will be holding the backup data of its own site, plus compressed copies of two other sites backups.

4.2 Replication management database

To allow all federated instances of CoStor to collaborate on managing replication tasks, and to enable restoration from federated data even if the master database for that site is

unavailable (in the event of a server failure), the replication data is synchronised across all servers.

All operations on these database tables are wrapped such that an API request is made to all other servers to update them immediately. Additionally, before a replication task commences, a full sync is requested to ensure that all servers contain all objects. To allow the servers to easily verify data integrity across nodes, all entries in this database have a UUID and are write only. When a sync is triggered, the server sends a request to all other CoStor servers with the full list of object IDs. If any IDs are missing, each server requests copies of these. If any server finds that it has an object with an ID that isn't in the list, it can assume that object has been deleted and will purge it from its own copy of the database.

This database contains information such as which backup Objects and file "primes" (Figure 3.1) have been replicated to which locations, current status of all nodes, and basic information on which sites own which agents and snapshots. This database **must** contain all the information needed to initiate a restore in the event of a remote disaster recovery restoration (Section 3.6). As only the server to be modifying data relating to a site is the server located in that site, there is no risk of requests interfering with each other.

As almost every object in CoStor is referenced by a unique and anonymous ID, there is minimal risk in sharing this data with other servers.

4.3 Securely synchronising data

The following concept relies heavily on file "primes" being stored in CoStor with convergent encryption[1], for cross-site deduplication, which has not yet been implemented in the core CoStor client and server applications.

"How do we share data across sites without risking breach of confidentiality?"

To enable the backup databases to be replicated securely, replicated data is stored using a modified subset of the standard database objects used by the local backups (Figure 3.1).

The actual DbFile objects do not require changing, as the data is already stored encrypted, and thanks to the use of convergent encryption, the ID (or hash) of all of these objects is already the same as those stored from the target server's local backups. The Object class is modified such that the name, path and stat fields are stored as encrypted strings, but as the IDs, agent and snapshot IDs are unique across sites, as well as being anonymous, these foreign keys can continue to be used.

4.4 Efficiently transferring data

We also need a way to transfer that data between hosts. As we already have serialisers for the database models thanks to the Django Rest Framework, it makes sense to re-use these to transfer database objects as JSON data. This also allows us to combine multiple entries into one large request, as an entire snapshot's worth of data is still fairly small.

For file primes, if they are not already present on the target server, we can give these a quick bit of compression using something fast like LZ4¹ to compress each object before transmitting it using a modified version of the large file protocol utilised by the client, described in Figure 2.1. These can then be decompressed and added to that site's local DbFile store, where they may come in useful for local backups as well.

¹<https://github.com/lz4/lz4>

Chapter 5

Testing and evaluation

Beyond being a cheap, self-hosted option to cloud backup, CoStor has two main goals:

5.1 Being a reliable backup solution

Testing for actual backup performance, and fault-tolerance, using a combination of real-world data such as test installs of operating systems, and backups of my various machines!

5.2 Being easy to use and manage

Testing the web management panel and installation process using user studies.

Bibliography

- [1] P. Anderson and L. Zhang, “Fast and secure laptop backups with encrypted de-duplication,” in *Proceedings of the Large Installations Systems Administration (LISA) Conference*. Berkeley, CA: Usenix Association, November 2010. [Online]. Available: <http://homepages.inf.ed.ac.uk/dcspaul/publications/lisa2010.pdf>
- [2] Syncthing website. [Online]. Available: <https://syncthing.net>
- [3] Urbackup website. [Online]. Available: <https://www.urbackup.org/index.html>
- [4] (2020, January) Hermes project on github. [Online]. Available: <https://github.com/Hermes/hermes>